

The 7-Layer CX Orchestration Framework

How the layers of a contact center stack connect, who usually owns each one, and what to check before you change any of them.

Layer 7: Measure and govern (Analytics + Governance)

Layer 6: Route the work (Routing + Orchestration)

Layer 5: Hold the conversation (Conversation Management)

Layer 4: Decide the next step (Reasoning + Planning)

Layer 3: Set the rules (Policy + Guardrails)

Layer 2: Do the work (Workflow Execution)

Layer 1: Know the customer (Data Access)

2026 edition, updated 29 September 2026. Replaces the April 2026 edition.

Independent. No vendor sponsorship. contactcentercx.com

Why a stack needs layers

A contact center runs on more than one platform. Customer data, workflows, rules, automated reasoning, the conversation itself, routing and measurement each have their own systems, owners and pace of change. When they are bought and run as one decision, the weakest connection sets the experience the customer gets.

This framework separates the stack into seven layers so each can be owned, tested and changed on its own terms. It gives you a shared vocabulary for the people who run the operation, the people who build the technology and the people who fund it.

Use it in four steps: place each of your systems on a layer; name who owns each layer today; read the questions for the layers you are about to change; then rate the readiness statements at the end and start with the weakest layer.

The seven layers at a glance

LAYER	WHAT IT DOES	WHO USUALLY OWNS IT
1. Know the customer	The systems of record the rest of the stack reads and writes: CRM, case management, billing, order and account systems, and the event streams between them.	Data and platform owners, with the business owner of each system of record.
2. Do the work	The workflows that carry out a request across systems: process engines, integration platforms and robotic automation.	Automation and integration engineers, with the operations owner of each process.
3. Set the rules	The rules every interaction must respect: payment security, identity verification, fraud controls, privacy and the guardrails on automated decisions.	Compliance, security, fraud and risk teams, with legal review of contracts.
4. Decide the next step	The reasoning that decides the next step: virtual agents, agent assist, knowledge retrieval and automated agents that act on a customer's behalf.	Conversation designers, knowledge owners and the CX architect, with risk review for anything that acts on its own.
5. Hold the conversation	The conversation itself: voice, chat, messaging, email, social and in-app support, and the agent workspace where they meet.	Channel and telephony owners, CX designers and the operations leaders who staff each channel.
6. Route the work	Routing and orchestration: which person or automated agent takes each contact, in what order, with what context and with what authority.	Routing and workforce specialists and the CX architect.

7. Measure and govern

Measurement and governance: quality management, workforce management, speech and text analytics, journey analytics and compliance monitoring.

Workforce, quality and analytics leads, with compliance.

The layers one by one

LAYER 1 · DATA ACCESS

Know the customer

The systems of record the rest of the stack reads and writes: CRM, case management, billing, order and account systems, and the event streams between them.

WHO USUALLY OWNS IT

Data and platform owners, with the business owner of each system of record.

QUESTIONS TO ANSWER

- Can the contact center read and write the customer record while the customer is still on the line, or does it wait for a later sync?
- Which systems does a typical contact touch, and who owns each one?
- Are interaction events (contact started, intent recognised, escalation) published somewhere the other layers can use them?

WHAT TO CHECK FIRST

List every system an agent or automated flow reads during a contact, with who owns it, how fresh its data is and whether it can be written back to. The gaps agents complain about most are usually the first integrations to fix.

WHEN IT FAILS

Every layer above loses context: automated answers cannot personalise, routing cannot prioritise, and agents ask customers for what the company already knows.

TEST IT WITH YOUR OWN NUMBERS

[CX IT Alignment](#)

Do the work

The workflows that carry out a request across systems: process engines, integration platforms and robotic automation.

WHO USUALLY OWNS IT

Automation and integration engineers, with the operations owner of each process.

QUESTIONS TO ANSWER

- Which of your highest-volume requests can be completed end to end without someone copying data between systems?
- When a workflow fails part way (a declined payment, a timeout, a policy exception), does it recover, or does the customer start again?
- Can your workflows be started through an API by an automated agent as well as by a person?

WHAT TO CHECK FIRST

Map your highest-volume workflows from start to finish: every system, every handoff and every point where one fails. Mark each as automated, partly automated or manual.

WHEN IT FAILS

The request is understood but not fulfilled. Agents complete it by hand and customers wait longer.

TEST IT WITH YOUR OWN NUMBERS

[TCO Calculator](#), [AHT Decomposition](#)

Set the rules

The rules every interaction must respect: payment security, identity verification, fraud controls, privacy and the guardrails on automated decisions.

WHO USUALLY OWNS IT

Compliance, security, fraud and risk teams, with legal review of contracts.

QUESTIONS TO ANSWER

- Do agents ever hear or type card numbers, or is payment handled outside the conversation?
- Is a customer verified the same way on voice, chat and self-service?
- For every point where an automated system answers a customer, what stops it promising what the business cannot deliver or disclosing what it should not?

WHAT TO CHECK FIRST

Document every point where an automated system makes or shapes a customer-facing decision, and the guardrail on each. If you serve customers in the EU, check which obligations of the AI Act apply to you and when, as amended by Regulation (EU) 2026/1744.

WHEN IT FAILS

Automated answers go ungoverned, payment flows fail and compliance exposure grows.

TEST IT WITH YOUR OWN NUMBERS

[Contract Risk](#), [Governance Model](#)

Decide the next step

The reasoning that decides the next step: virtual agents, agent assist, knowledge retrieval and automated agents that act on a customer's behalf.

WHO USUALLY OWNS IT

Conversation designers, knowledge owners and the CX architect, with risk review for anything that acts on its own.

QUESTIONS TO ANSWER

- Does your virtual agent answer from your current knowledge, policies and product data, or only from what it was built with?
- Does agent assist help during the conversation, or only summarise after it ends?
- What share of contacts does automation fully resolve today, measured the same way a customer would judge it?

WHAT TO CHECK FIRST

Review your knowledge base on a schedule you set and keep: who owns each article, when it was last checked, and what the automated layer does when the answer is not there.

WHEN IT FAILS

Self-service fails, contacts that could have been resolved flow to agents, and queues grow.

PUBLISHED FIGURES: SELF-SERVICE AND AI RESOLUTION

14% Issues fully resolved in self-service

Gartner, 19 Aug 2024. [Source](#)

18% Web chats handled without a human agent, US

ContactBabel, 2026. [Source](#)

30% Service cases handled by AI

Salesforce, 13 Nov 2025. [Source](#)

TEST IT WITH YOUR OWN NUMBERS

[AI Deflection Reality Check](#), [AI Readiness](#)

Hold the conversation

The conversation itself: voice, chat, messaging, email, social and in-app support, and the agent workspace where they meet.

WHO USUALLY OWNS IT

Channel and telephony owners, CX designers and the operations leaders who staff each channel.

QUESTIONS TO ANSWER

- If a customer moves from chat to a call, does the agent see what was already said?
- Can your digital channels carry a complex conversation to resolution, or only simple questions?
- How many applications does an agent open to resolve a typical contact?

WHAT TO CHECK FIRST

Test continuity yourself: start a conversation in one channel, move it to another, and see what carries over. Count the applications an agent opens on a typical contact.

WHEN IT FAILS

Customers repeat themselves, channels fall out of step, and the measurement layer loses the thread.

PUBLISHED FIGURES: CHANNEL MIX AND PREFERENCE

62% Live telephony, share of inbound, end of 2024
ContactBabel, read on Contact Center Pipeline, 14 Aug 2025. [Source](#)

35% Phone as the preferred channel, US adults
YouGov, 13 Mar 2025. [Source](#)

TEST IT WITH YOUR OWN NUMBERS

[Channel Shift](#), [FCR Leakage](#)

Route the work

Routing and orchestration: which person or automated agent takes each contact, in what order, with what context and with what authority.

WHO USUALLY OWNS IT

Routing and workforce specialists and the CX architect.

QUESTIONS TO ANSWER

- Do you route on skill tags alone, or also on intent, customer history and the capacity actually available?
- Can routing send a contact to an automated agent first and hand it to a person, with context, when needed?
- Does routing coordinate with workforce management, the CRM and knowledge, or run on its own?

WHAT TO CHECK FIRST

List your routing rules and when each was last reviewed, and check that the staffing plan behind them matches the demand the rules create.

WHEN IT FAILS

Contacts reach the wrong person or the wrong automated agent, transfers rise and customer effort climbs.

PUBLISHED FIGURES: SERVICE LEVEL TARGETS

80% in 20 s	Traditional service level standard SQM Group, 19 Jan 2023. Source
--------------------	--

80% in 30 s	California water utilities, Class A and B California Public Utilities Commission, 2009. Source
--------------------	---

2:00	Medicare Advantage lines, CMS standard CMS, 2025. Source
-------------	---

74 s	Average speed of answer, US ContactBabel, 2026. Source
-------------	---

6%	Abandon rate, industry standard SQM Group, 19 Jan 2023. Source
-----------	---

TEST IT WITH YOUR OWN NUMBERS

[Staffing Calculator](#), [Occupancy Risk](#)

Measure and govern

Measurement and governance: quality management, workforce management, speech and text analytics, journey analytics and compliance monitoring.

WHO USUALLY OWNS IT

Workforce, quality and analytics leads, with compliance.

QUESTIONS TO ANSWER

- How are interactions chosen for quality review, and does the sample size answer the questions you ask of it?
- How far are your forecasts from actual volume, measured interval by interval?
- Do findings from analytics change routing, coaching and automated answers, or stop at a dashboard?

WHAT TO CHECK FIRST

Measure your forecast accuracy by interval, calibrate your quality evaluators against each other, and trace one recent finding from the dashboard to the change it produced.

WHEN IT FAILS

Nothing stops at once, but quality drifts unseen, staffing falls out of step with demand and compliance gaps widen.

PUBLISHED FIGURES: SCHEDULE ADHERENCE

85% to 95% Generally considered acceptable
Call Centre Helper, updated 22 Jul 2026. [Source](#)

85% to 95% High-performing centers target, from a workforce software vendor
Verint, 1 Jan 2021. [Source](#)

80% What most centers strive for
TechTarget, 14 Jun 2024. [Source](#)

TEST IT WITH YOUR OWN NUMBERS

[Forecast Accuracy](#), [QA Scorecard](#), [Schedule Adherence](#)

What depends on what

No layer works alone. Most failed projects trace back to a dependency nobody mapped before the work began. This is our reading of how a failure travels up and across the stack.

IF THIS LAYER FAILS	THESE FEEL IT	WHAT THE CUSTOMER NOTICES
1. Know the customer	Every layer above	Agents and automated answers work without context.
2. Do the work	Layers 4, 5 and 6	Requests are understood but not completed.
3. Set the rules	Layers 4 and 5	Automated answers and payments run without the rules they need.
4. Decide the next step	Layers 5 and 6	Contacts that automation could resolve reach agents instead.
5. Hold the conversation	Layers 6 and 7	Customers cannot reach you in the channel they chose, and measurement loses data.
6. Route the work	Layers 5 and 7	Contacts reach the wrong resource and effort rises.
7. Measure and govern	No layer at once	Quality, staffing and compliance drift without anyone seeing it.

Start at the bottom. Automated reasoning built on incomplete or stale data gives confident answers that are wrong. Fixing the data layer first makes every change above it safer.

Readiness checklist

Rate each statement from 1 to 5: 1 not started, 2 partly in place, 3 working, 4 working and reviewed, 5 working, reviewed and improving.

Every statement you rate 2 or below is an action. Start with the layer that has the most of them, lowest layer first when two are level. There is no total and no band: a strong layer does not make up for a weak one.

LAYER	STATEMENT	YOUR RATING
1. Know the customer	The contact center can read and write customer, billing and case data while the customer is on the line.	1 2 3 4 5
	Every system the contact center depends on is documented with its owner.	1 2 3 4 5
2. Do the work	Our highest-volume workflows are documented end to end, with the systems each one touches.	1 2 3 4 5
	Our highest-volume requests complete without anyone copying data between systems.	1 2 3 4 5
3. Set the rules	Card payments are handled outside the conversation.	1 2 3 4 5
	Every point where automation answers a customer has a documented guardrail and an owner.	1 2 3 4 5
4. Decide the next step	Automated answers draw on current knowledge, policies and product data.	1 2 3 4 5
	The knowledge base is reviewed on a schedule we keep, and each article has an owner.	1 2 3 4 5
5. Hold the conversation	Context carries over when a customer moves between chat, voice and self-service.	1 2 3 4 5
	Agents resolve a typical contact from one workspace.	1 2 3 4 5
6. Route the work	Routing rules are reviewed on a schedule and reflect current priorities.	1 2 3 4 5
	Routing can send a contact to an automated agent first and hand it to a person with context.	1 2 3 4 5
7. Measure and govern	The quality sample is chosen on purpose and big enough for the questions we ask of it.	1 2 3 4 5

7. Measure and govern	Findings from analytics change routing, coaching and automated answers.	1 2 3 4 5
------------------------------	---	-----------

For a scored assessment with a published rubric, try [CX IT Alignment](#) or [CX Maturity](#).

About this framework

What it is. A planning framework: the decisions each layer holds, who usually holds them, and what to check before you change anything. It is our own synthesis, in our own words.

What it is not. It is not a vendor evaluation. It does not score, rank or name platforms, and it does not replace architecture design for your operation.

Figures. Figures appear only where a publisher reports one, each with its source and date. Everything else is a question to answer with your own numbers, and each layer names a free diagnostic on the site that does it.

What changed from the April 2026 edition

- Unsourced figures are removed, including a 2% quality sample, a 10% forecast gap and fixed review intervals.
- Predictions about the next 12 months are removed; each layer now says what to check today.
- The EU AI Act note now points to the obligations as amended by Regulation (EU) 2026/1744.
- The readiness checklist no longer totals to a maturity band; a statement rated 2 or below is an action.
- Links to retired tools are replaced with the diagnostics that exist today.

Related: [Platforms and Tech](#), [Platform Decision](#), [TCO Calculator](#), [Vendor profiles](#), [Get help](#).

Published by The Center of CX, 29 September 2026. Report an error at contactcentercx.com/corrections.